



Trinity College Dublin
Coláiste na Tríonóide, Baile Átha Cliath
The University of Dublin

03 – Stacks

CS1022 – Introduction to Computing II

Dr Adam Taylor / adam.taylor@tcd.ie
School of Computer Science and
Statistics

A stack is an example of an abstract data type

- A model for organising data

- Defined operations and behaviour

Fundamental to the operation of most computers

- (particularly the implementation of “methods” / “functions” / “procedures” / “subroutines”)

Convenient data structure for other tasks

Analogous to a stack of paper or a stack of cards

Operations

“Push”: Place item on the top of the stack

“Pop”: Remove item from the top of the stack


In practice, we can observe or change items anywhere in the stack

But this goes beyond the normal (formal) definition of a stack

A LIFO data structure: Last In First Out


Compare with FIFO: First In First Out
(guess what we call this ...)

See Algorithms and Data Structures next year!




Stack Example


4




Push 8




Push 14




Push 22




Pop




Pop



Push 66



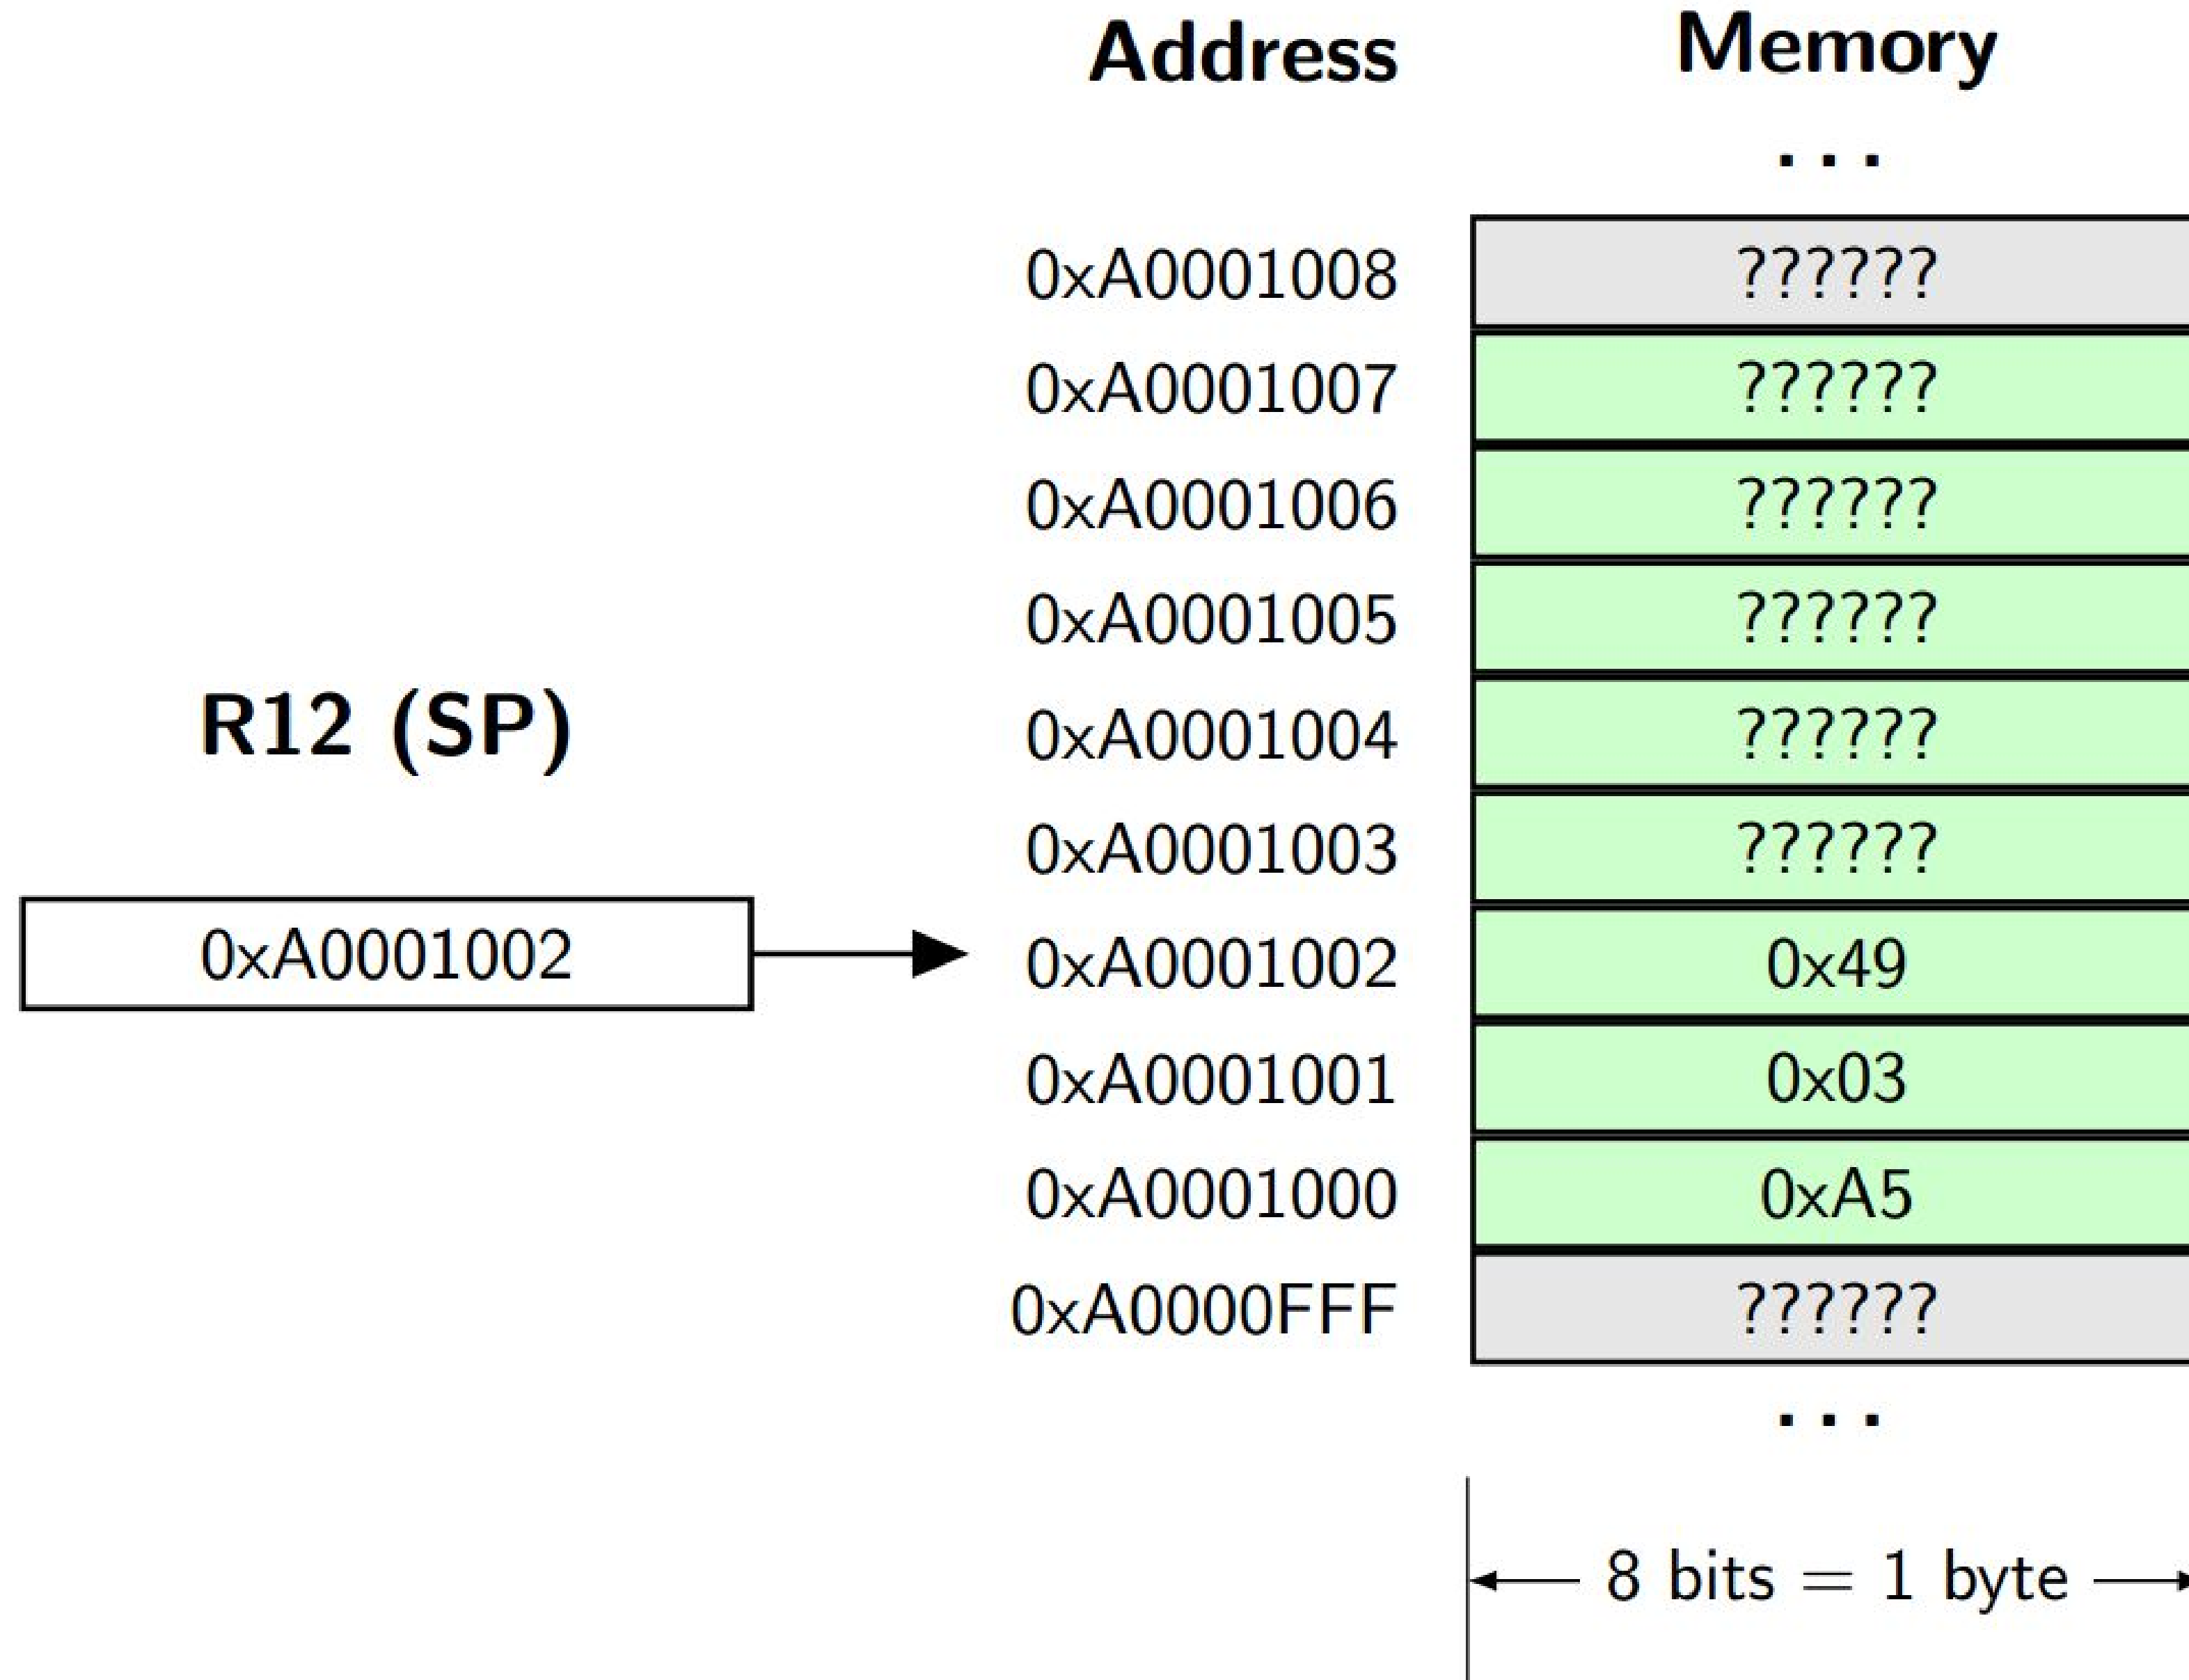
Pop, Pop



To implement a stack we need ...

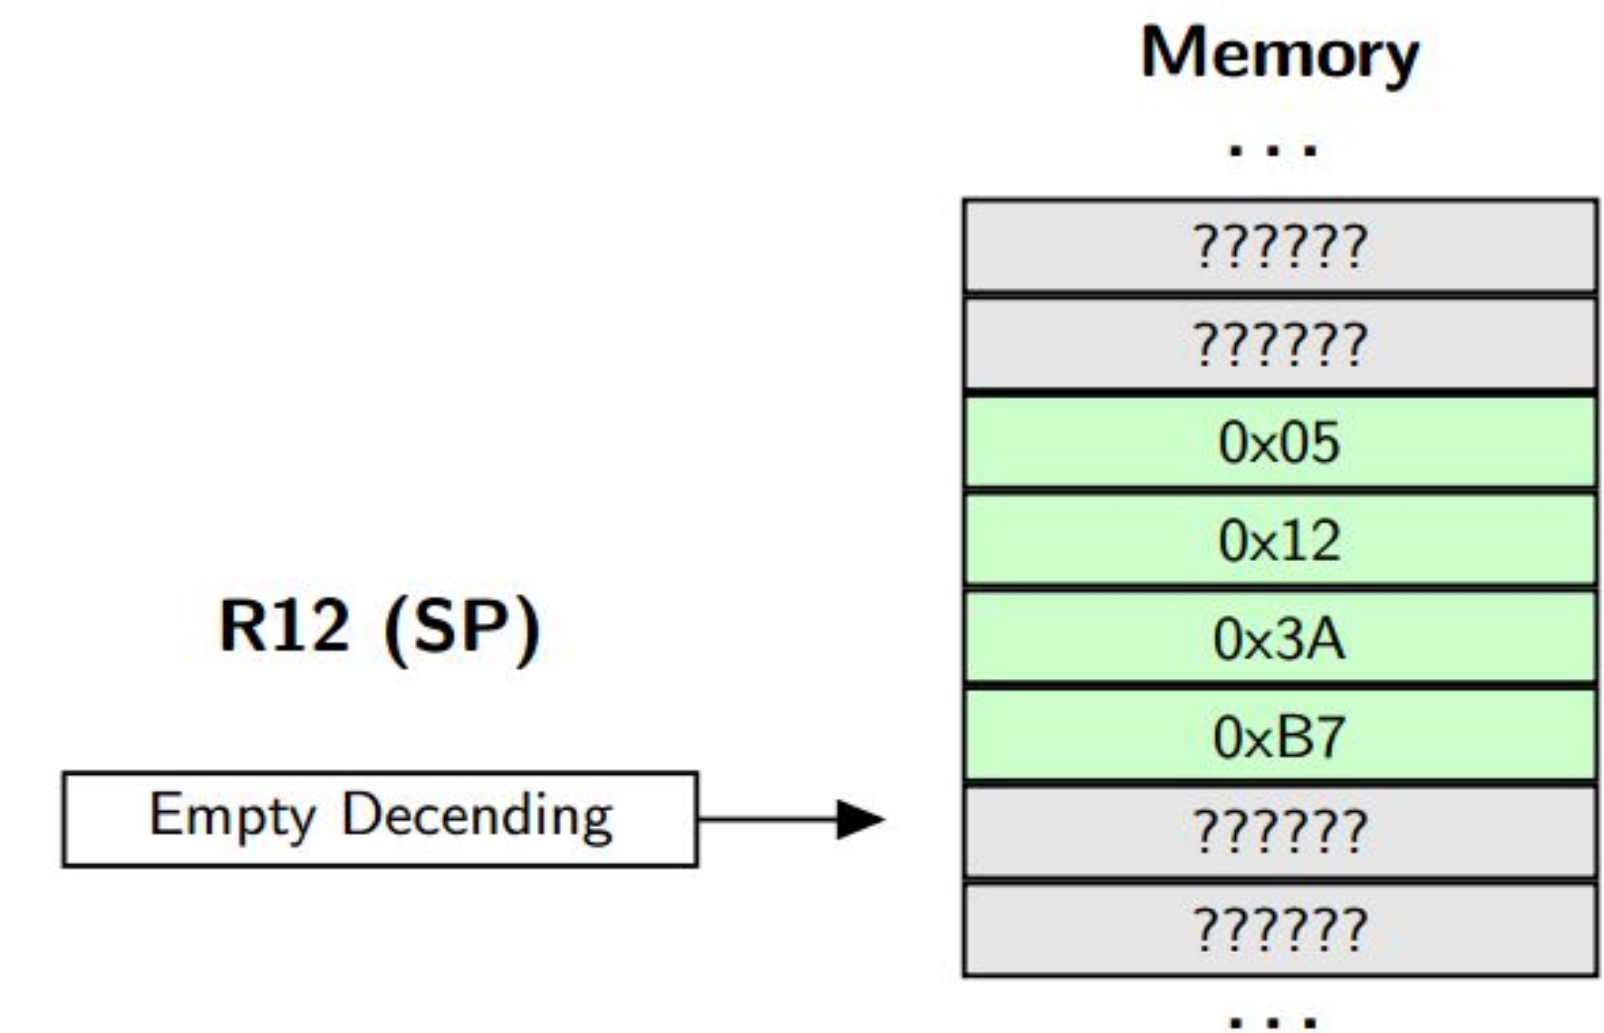
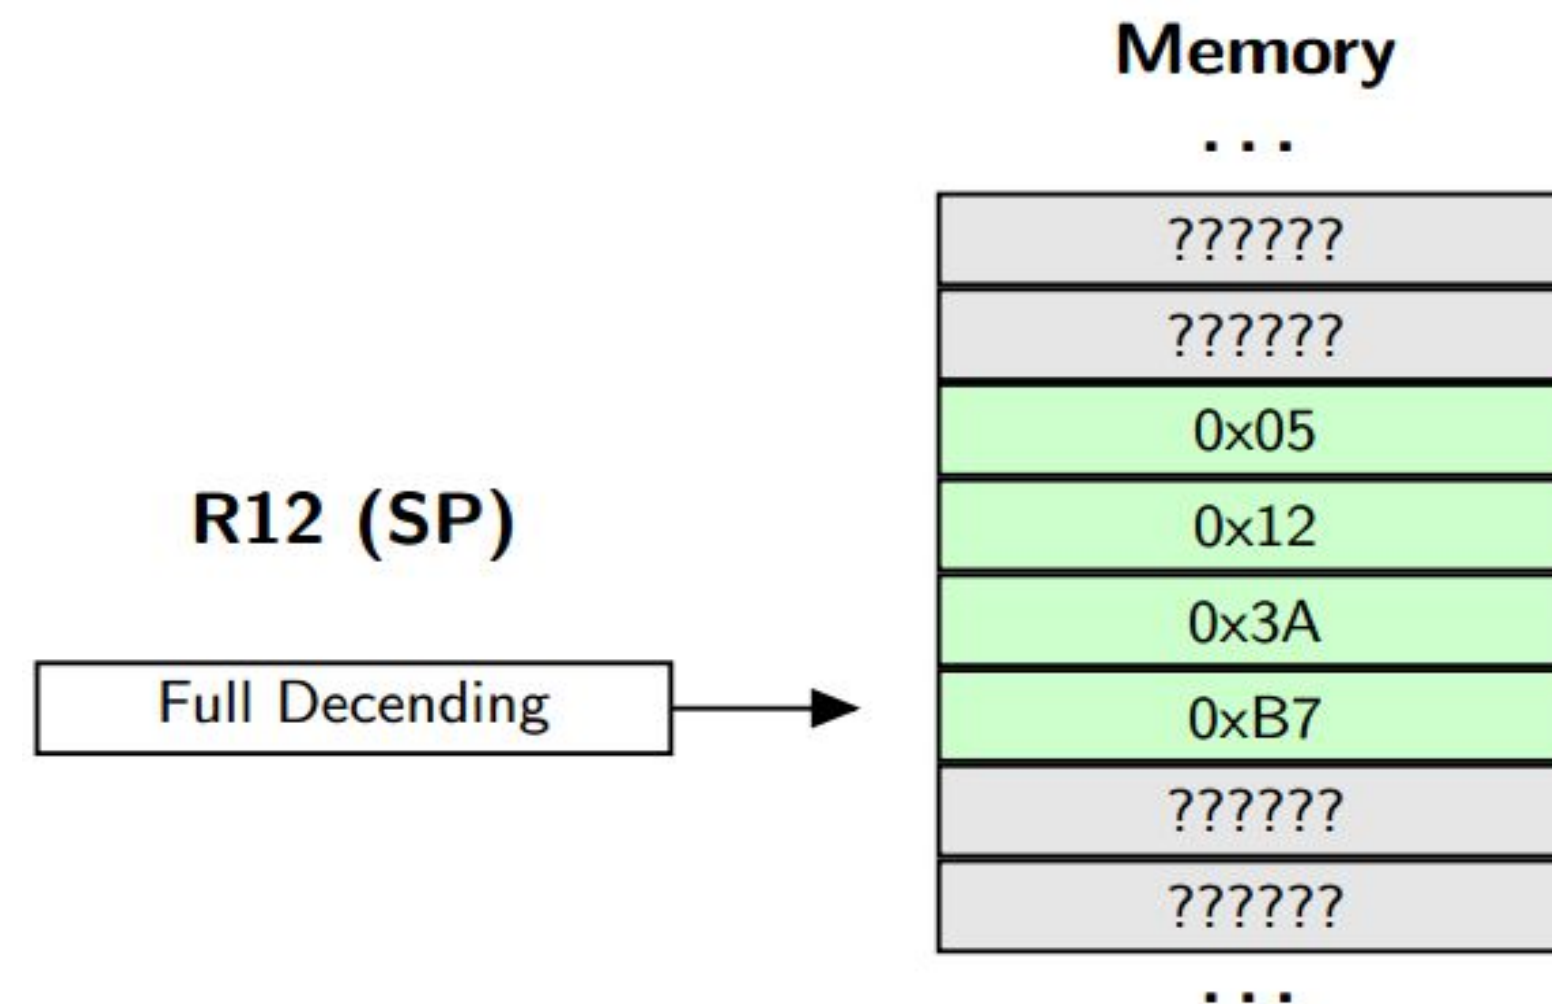
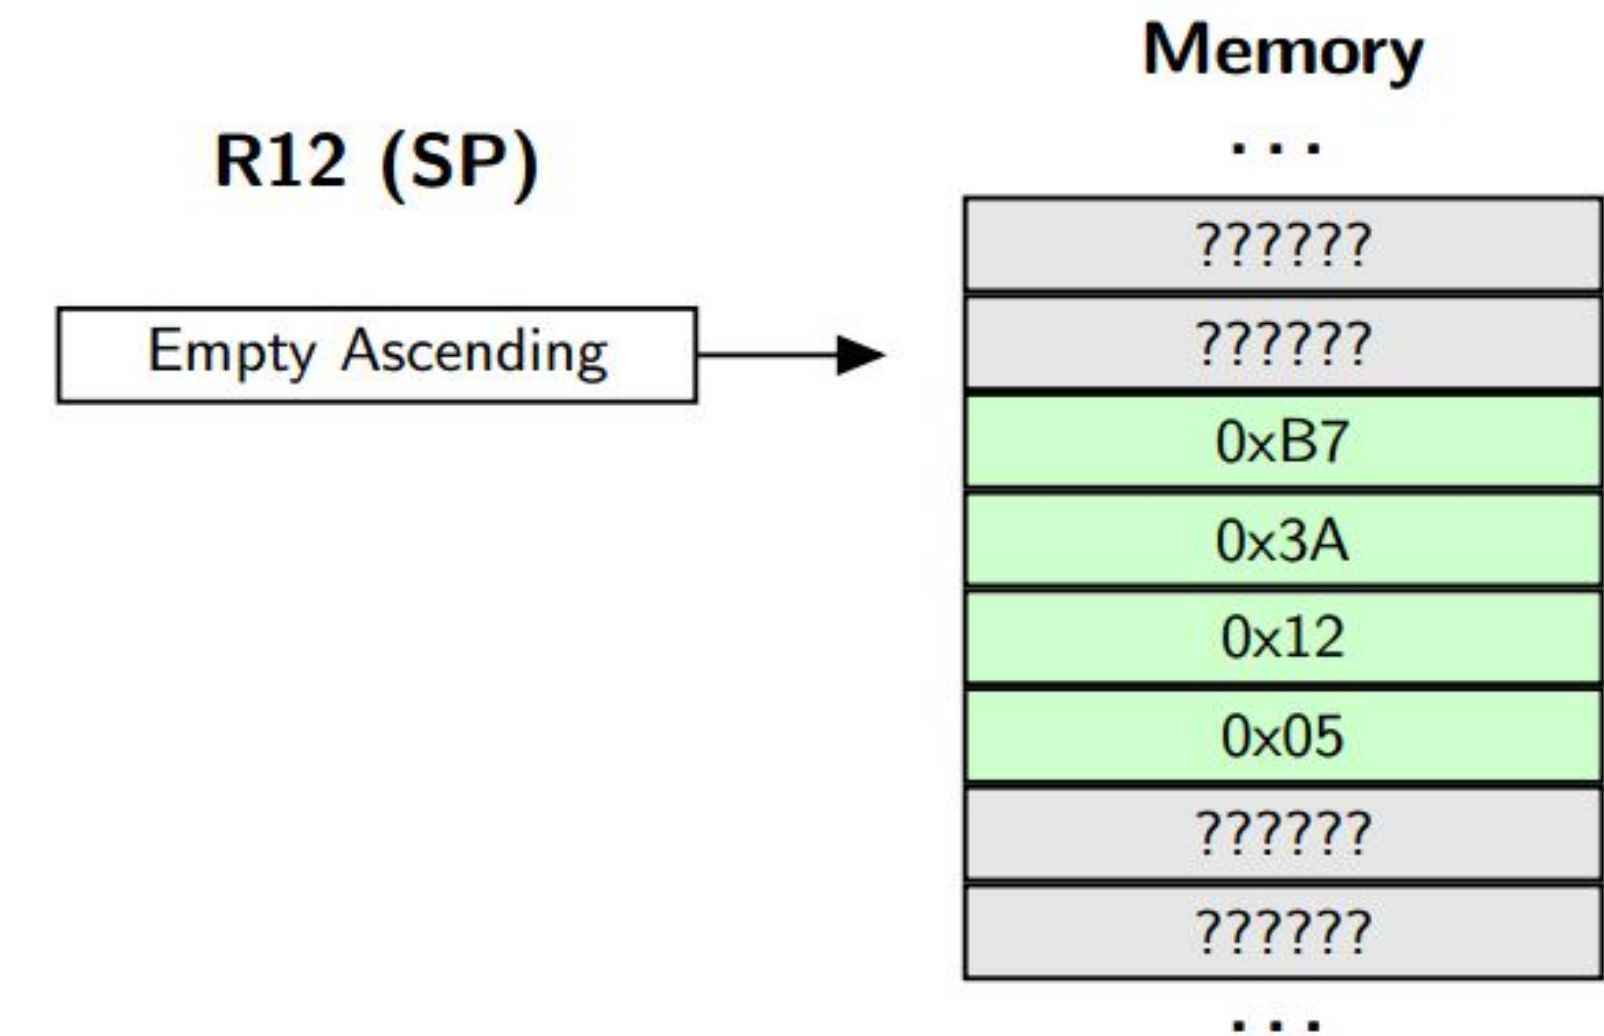
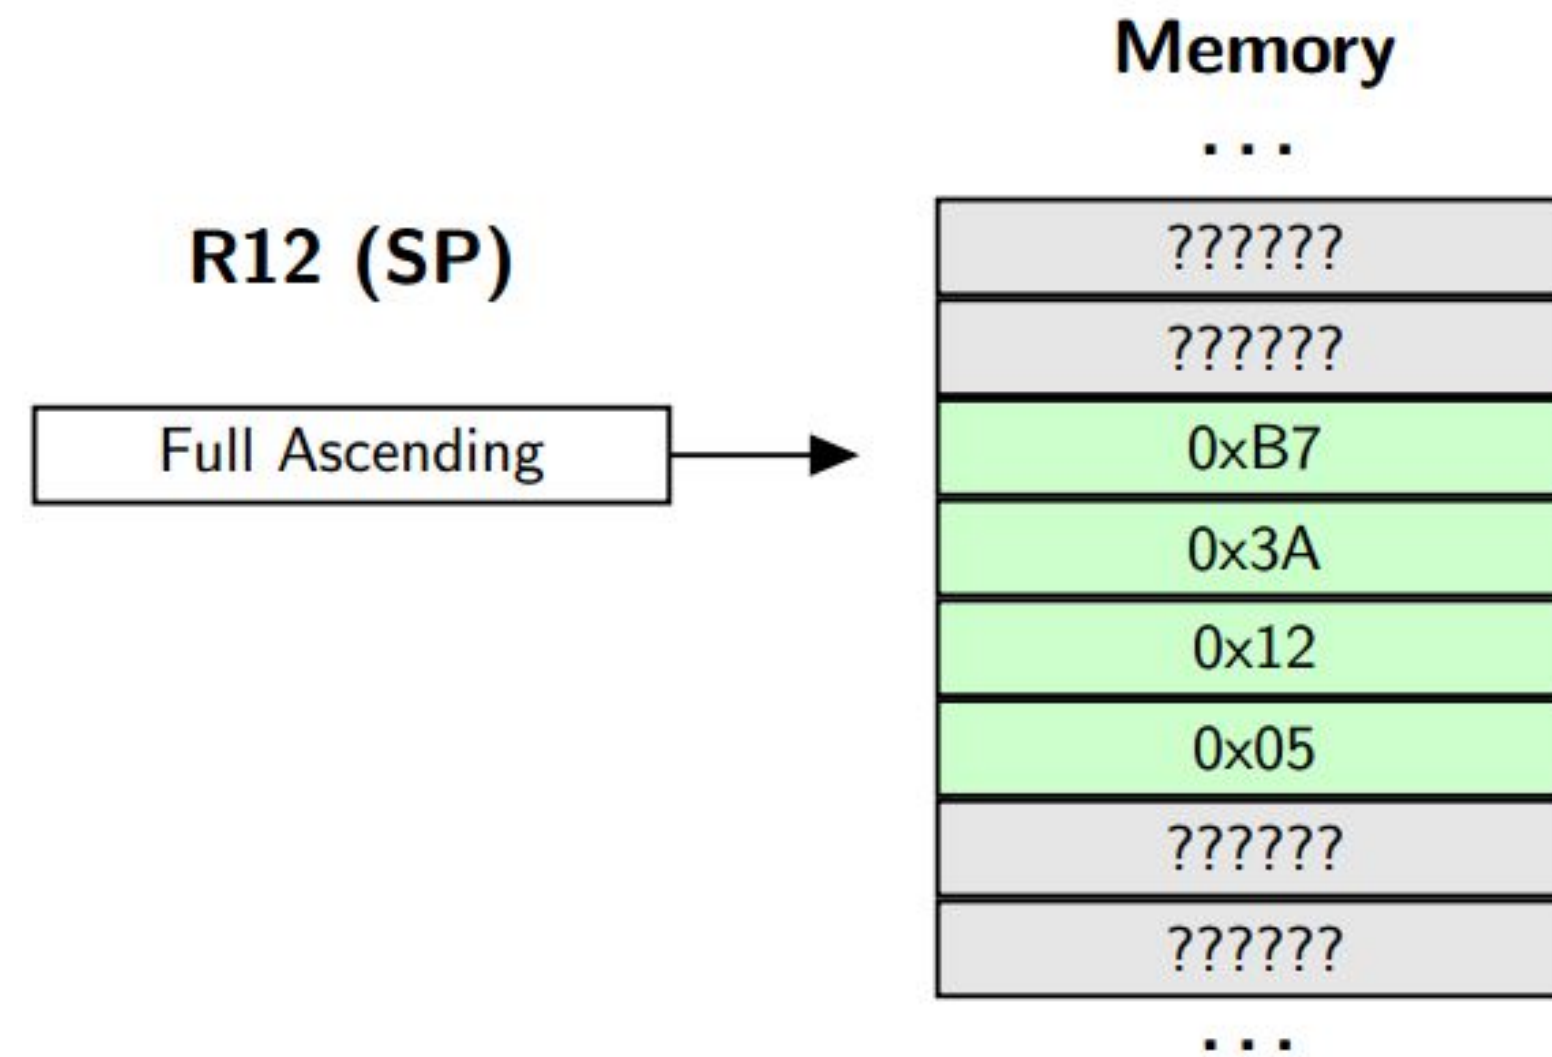

1. An area of memory to store the data items
size of the area of memory determines the maximum size of the stack
2. A Stack Pointer (SP) register to point to the top of the stack
we will see that we don't need to know where the bottom of the stack is!!
3. A stack growth convention (rules for pushing and popping)

Ascending or Descending	Full or Empty
Does the stack grow from low to high (ascending stack) or from high to low (descending stack) memory addresses?	Does the stack pointer point to the last item pushed onto the stack (full stack), or the next free space on the stack (empty stack)?



Stack Growth Convention

7



Initialisation

Set Stack Pointer (SP) to address at the start or end of the memory region to be used to store the stack (must consider the growth convention)

This is the bottom of the stack (and, since the stack has just been initialised, also the top of the stack!)

```
start
    LDR    R12, =STK_TOP      ; Initialise stack pointer
                                ; (assume full descending)
    ...    ...                ; ...
    ...    ...                ; rest of program
    ...    ...                ; ...

stop B    stop

STK_SZ EQU 0x400             ; 1K stack

    AREA Stack, DATA, READWRITE
    STK_MEM SPACE STK_SZ
    STK_TOP
```


Stack Implementation – push

9

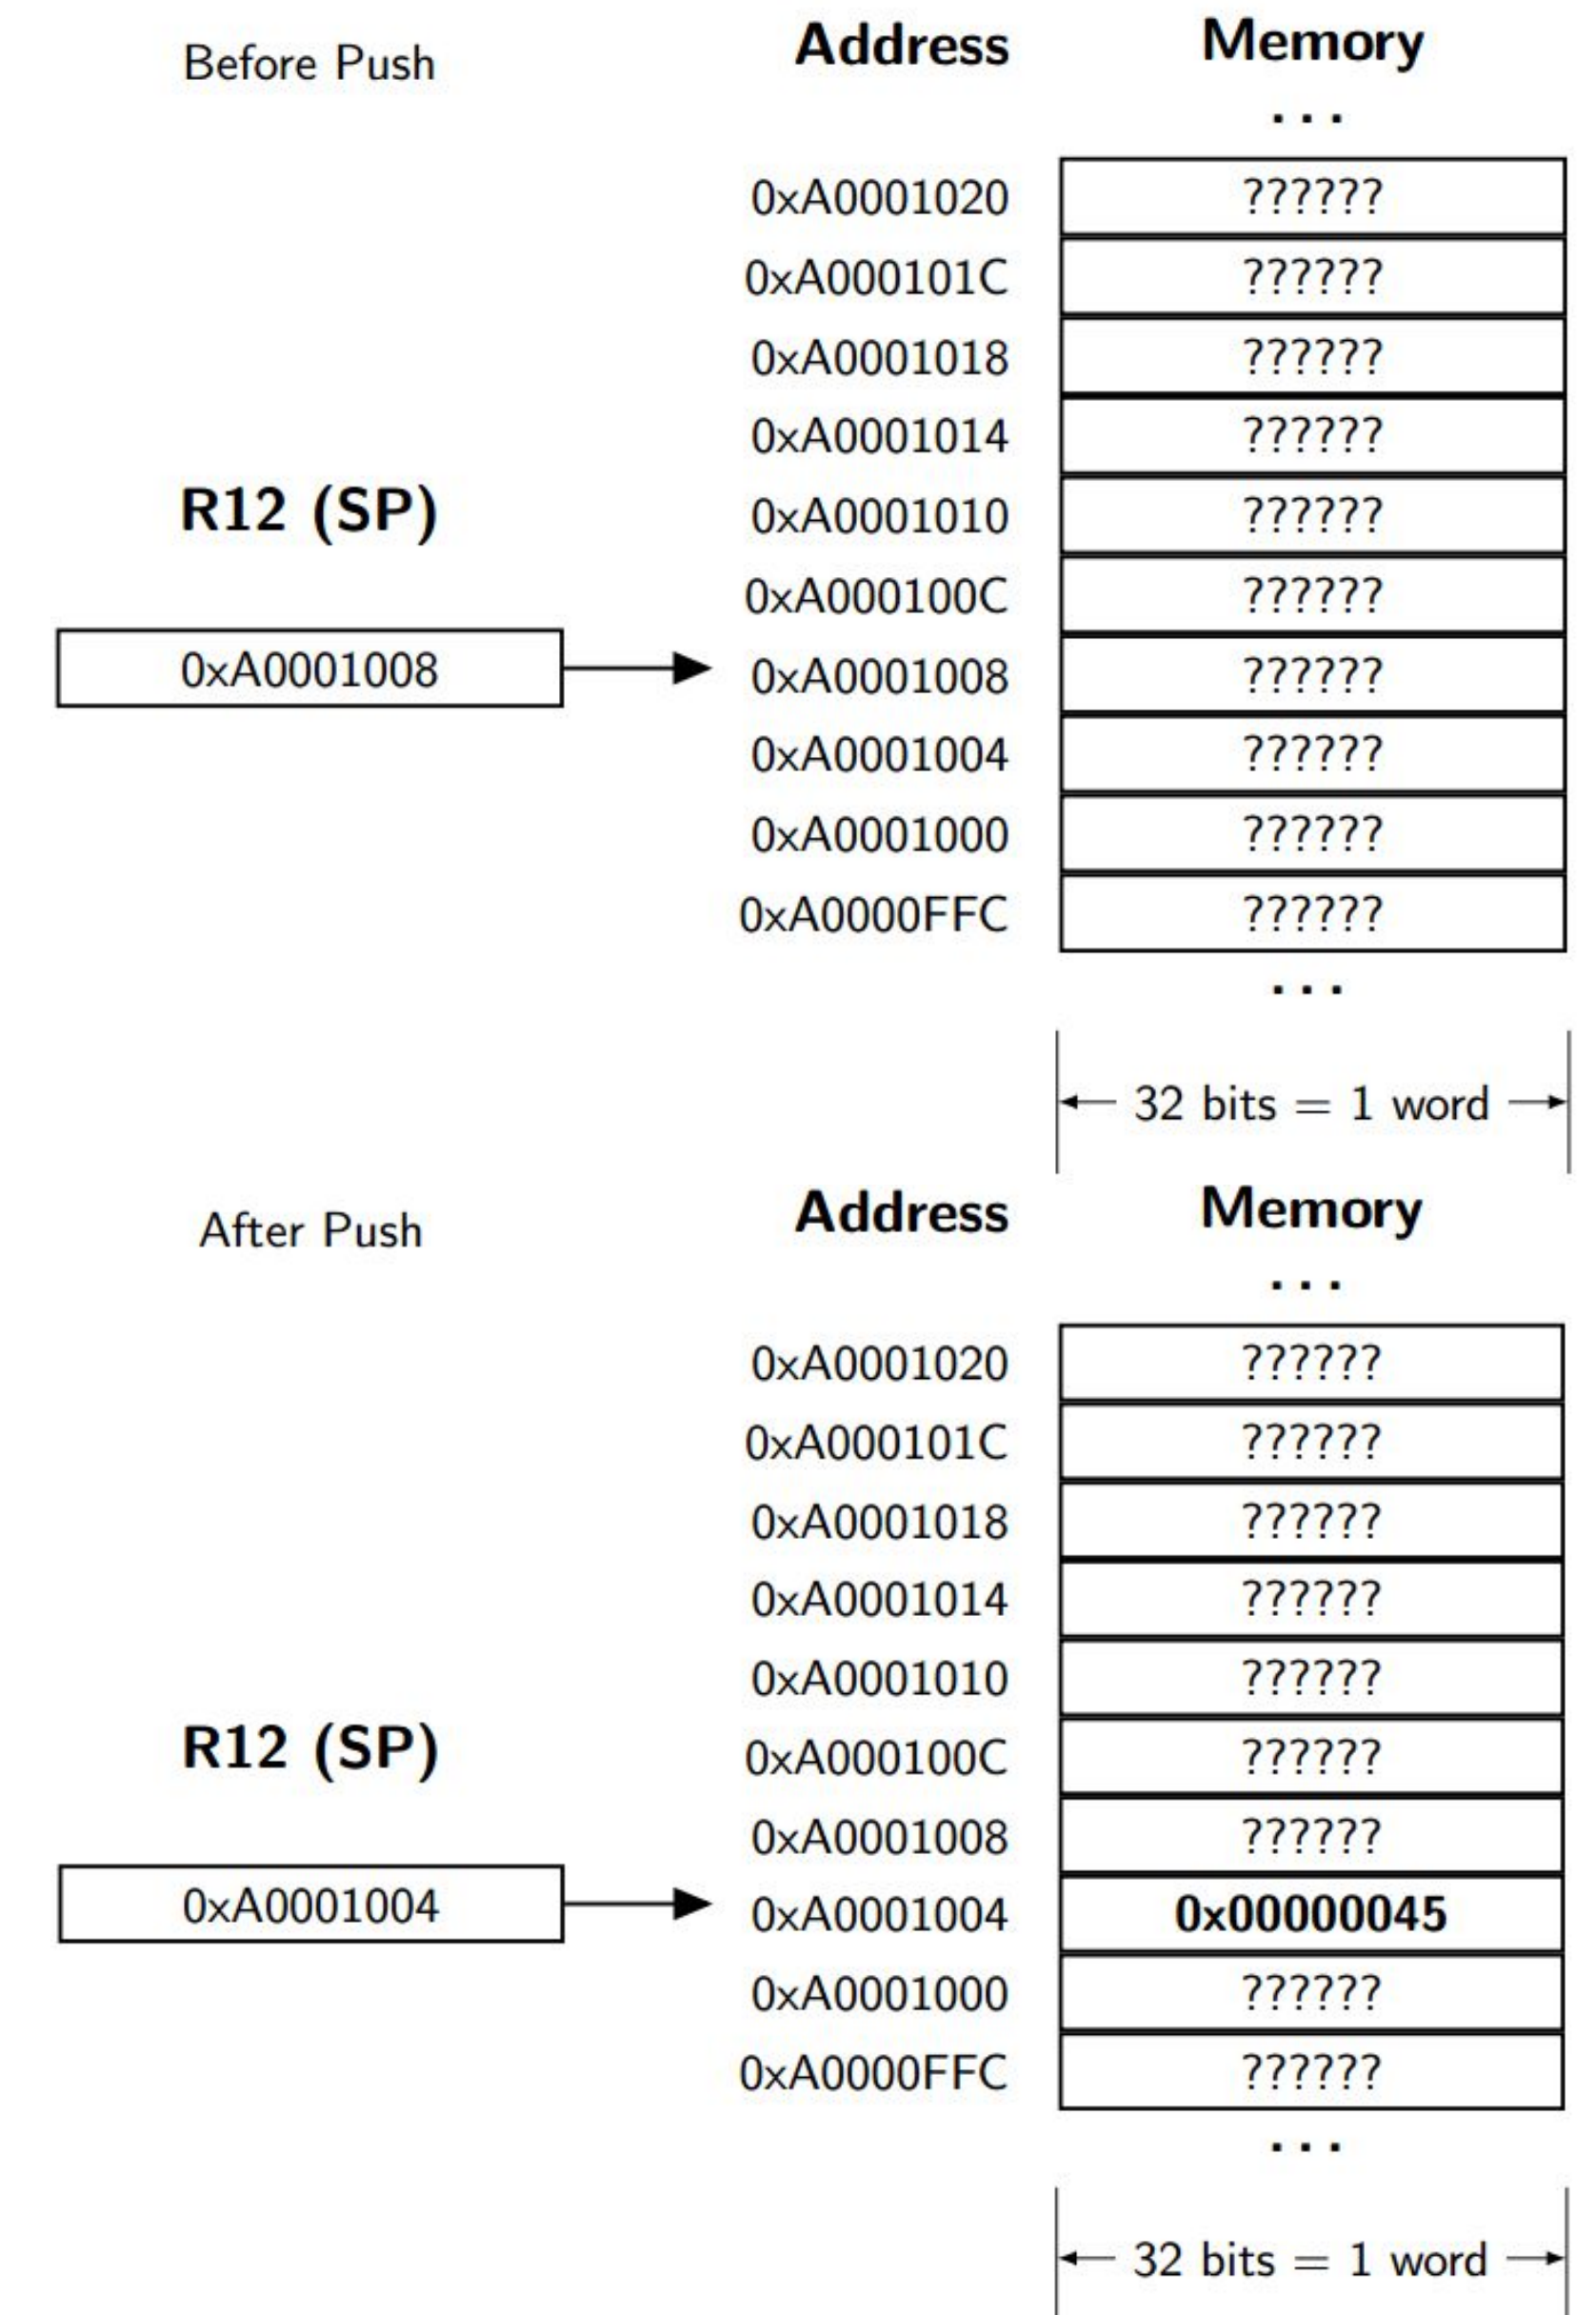
Assume a full descending stack growth convention

To push a word onto the stack

1. decrement the stack pointer by 4 bytes
(4 bytes = 1 word = 32 bits)
2. store the word in memory at the location pointed to by the stack pointer

e.g. push 0x45 using R12 as stack pointer

```
LDR R0, =0x45 ; test value
SUB R12, R12, #4 ; adjust SP
STR R0, [R12]
```



e.g. Push three words

```
; push 0x45  
LDR R0, =0x45  
SUB R12, R12, #4  
STR R0, [R12]
```

```
; push 0x7b  
LDR R0, =0x7b  
SUB R12, R12, #4  
STR R0, [R12]
```

```
; push 0x19  
LDR R0, =0x19  
SUB R12, R12, #4  
STR R0, [R12]
```

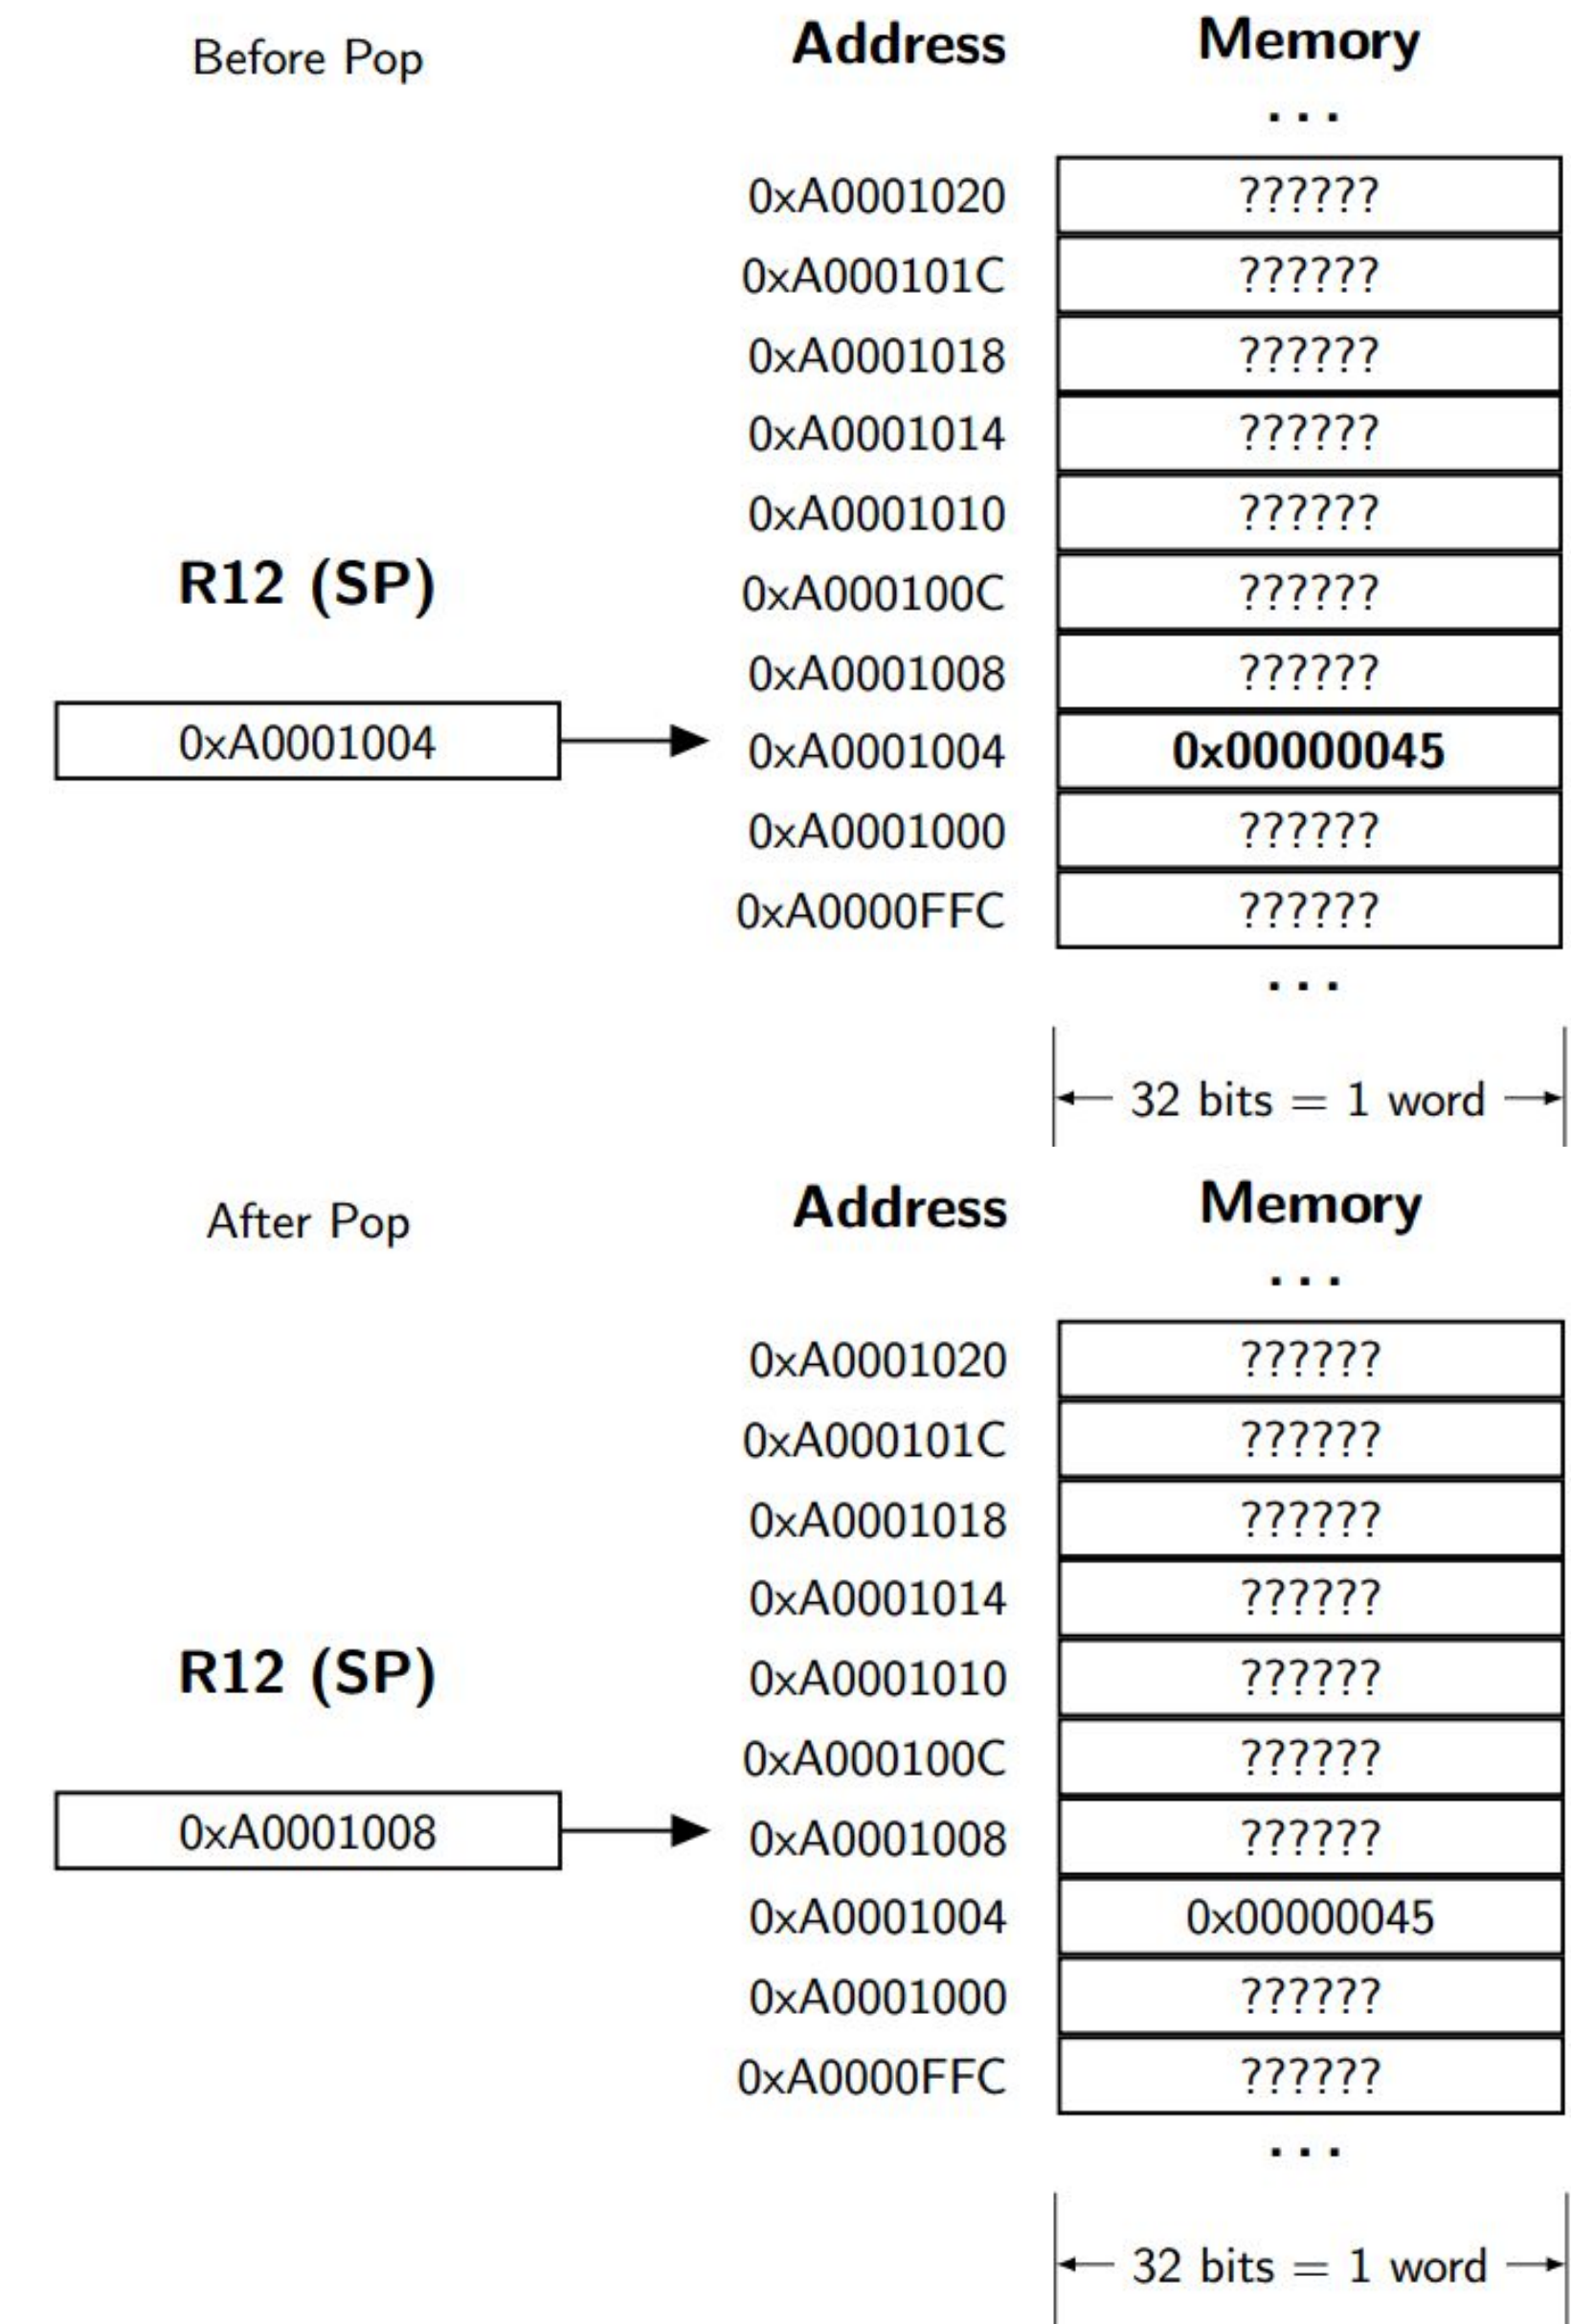
Again, assume full descending stack growth convention

To pop a word off the stack

1. load the word from memory at the location pointed to by the stack pointer (into a register)
2. increment the stack pointer by 4 bytes

e.g. pop word off top of stack into R0

```
LDR R0, [R12]
ADD R12, R12, #4
```



e.g. Pop three word-size values off the top of the stack

```
; pop  
LDR R0, [R12]  
ADD R12, R12, #4
```

```
; pop  
LDR R0, [R12]  
ADD R12, R12, #4
```

```
; pop  
LDR R0, [R12]  
ADD R12, R12, #4
```

Contents of R0 after each pop operation depend on contents of stack

e.g. if we had previously pushed 0x45, 0x7b and 0x19, we will pop 0x19, 0x7b and 0x45

e.g. Push word from R0 to stack pointed to by R12

```
; push word from R0
```

```
SUB R12, R12, #4
```

```
STR R0, [R12]
```

Replace explicit SUB with immediate pre-indexed addressing mode

```
; push word from R0
```

```
STR R0, [R12, #-4]!
```

Important!

Similarly, to pop word, replace explicit ADD with immediate post-indexed addressing mode

```
; pop word into R0
```

```
LDR R0, [R12], #4
```

Important!

In general, stacks ...

- can be located anywhere in memory

- can use any register as the stack pointer

- can grow as long as there is space in memory

Usually, a computer system will provide one or more system-wide stacks to implement certain behaviour (in particular, subroutine calls)

- ARM processors use register R13 as the system stack pointer (SP)

- Stack pointer is initialised by startup code executed when the computer is powered-on

- (libcs1021.lib contains our startup code)

- Limited in size (stack overflow)

Rarely any need to use any other stack

Use the system stack pointed to by R13/SP for your own purposes

```
; push word from R0  
STR R0, [SP, #-4]!
```

Note use of SP in place of R13

Never re-initialise R13/SP

```
; load address 0xA1000000 into R13  
LDR R13, =0xA1000000
```

Don't do something like this!!

Typical use of a system stack is temporary storage of register contents

Programmer's responsibility to pop off everything that was pushed on to the system stack

Not doing this is likely to result in an error that may be very hard to find!!

Frequently we need to load/store the contents of a number of registers from/to memory

```
; store contents of R1, R2 and R3 to memory at address 0xA1001000
```

```
LDR R12, =0xA1001000 ; initialise R12 with base address
```

```
STR R1, [R12]
```

```
STR R2, [R12, #4]
```

```
STR R3, [R12, #8]
```

```
; load R1, R2 and R3 with contents of memory at address 0xA1001000
```

```
LDR R12, =0xA1001000 ; initialise R12 with base address
```

```
LDR R1, [R12]
```

```
LDR R2, [R12, #4]
```

```
LDR R3, [R12, #8]
```

WARNING!



This is not a stack

we're leaving stacks for
a moment ...

ARM instruction set provides LoaD Multiple (LDM) and STore Multiple (STM) instructions for this purpose

The following examples achieve the same end result as the previous example ...

```
; store contents of R1, R2 and R3 to memory at address 0xA1001000
```

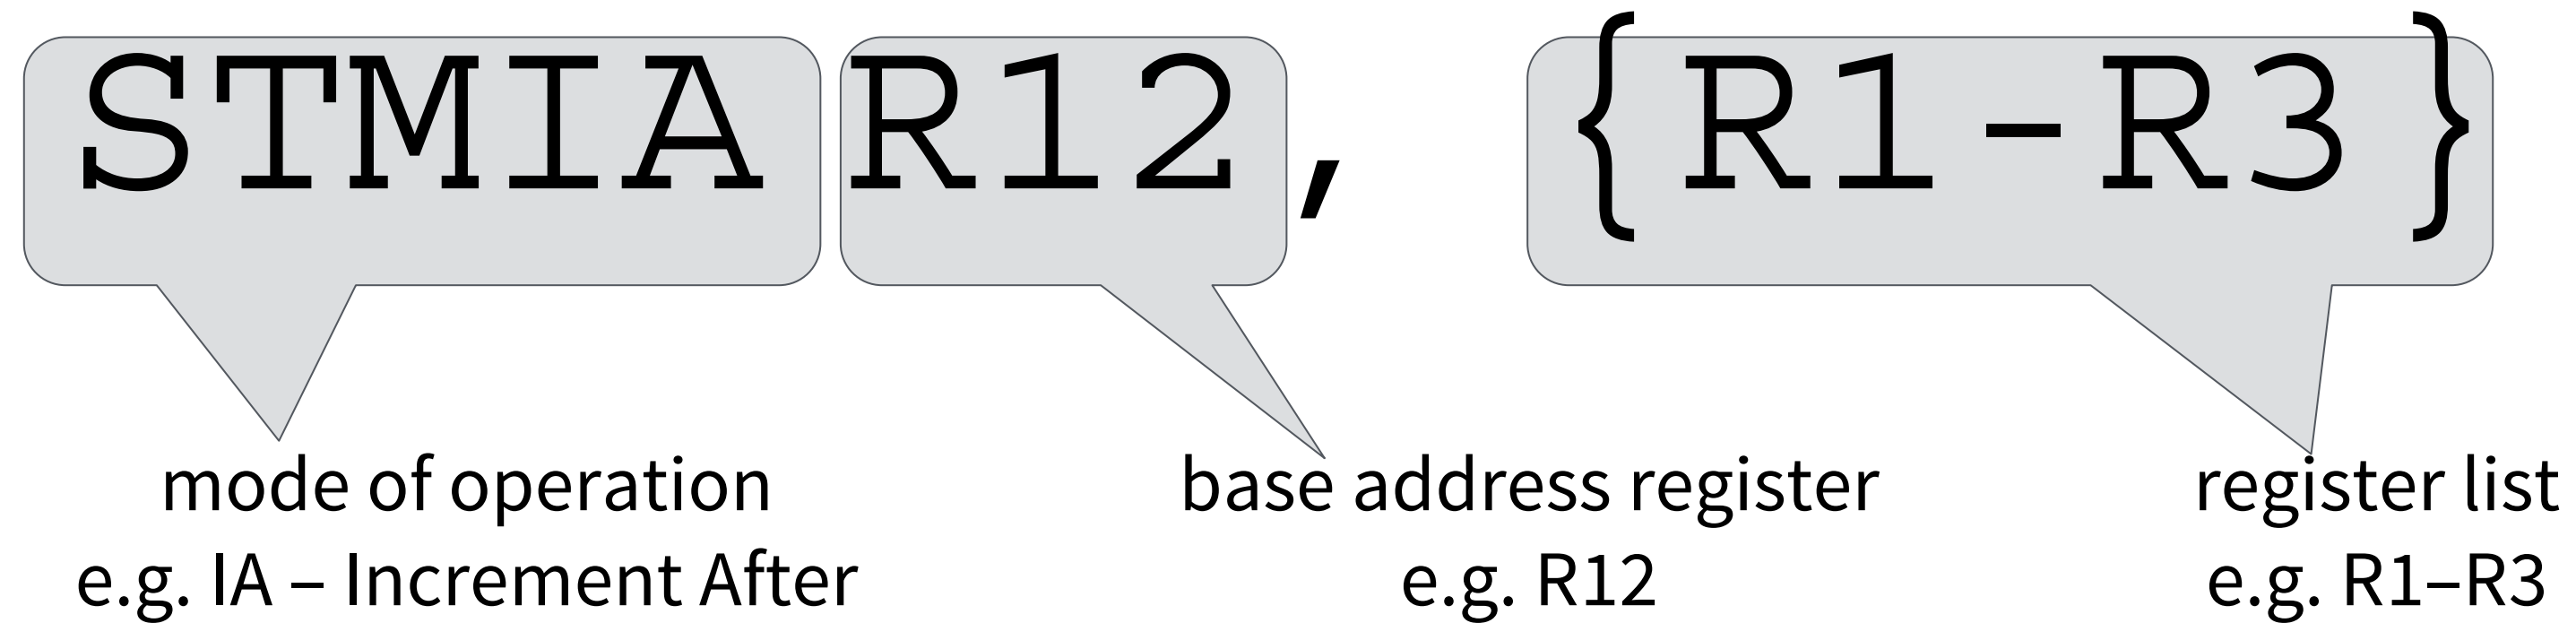
```
LDR R12, =0xA1001000  
STMIA R12, {R1-R3}
```

```
; load R1, R2 and R3 with contents of memory at address 0xA1001000
```

```
LDR R12, =0xA1001000  
LDMIA R12, {R1-R3}
```



Consider the following STM instruction ...



Note position of base address register operand!!

Increment After (IA) mode of operation:

first register is stored at <base address>

second register is stored at <base address> + 4

third register is stored at <base address> + 8

Contents of base register R12 remain unchanged



STMIA R12, {R1-R3}

Before STMIA

Address

Memory

...

0xA0001020

??????

0xA000101C

??????

0xA0001018

??????

0xA0001014

??????

0xA0001010

??????

0xA000100C

??????

0xA0001008

0xA0001008

??????

0xA0001004

??????

0xA0001000

??????

0xA0000FFC

??????

...

← 32 bits = 1 word →

After STMIA

Address

Memory

...

0xA0001020

??????

0xA000101C

??????

0xA0001018

??????

0xA0001014

??????

0xA0001010

R3

0xA000100C

R2

0xA0001008

0xA0001008

R1

0xA0001004

??????

0xA0001000

??????

0xA0000FFC

??????

...

← 32 bits = 1 word →



Four modes of operation for LDM and STM instructions

Behaviour	LDM	STM
Increment After	LDMIA	STMIA
Increment Before	LDMIB	STMIB
Decrement After	LDMDA	STMDA
Decrement Before	LDMDB	STMDB

Register list

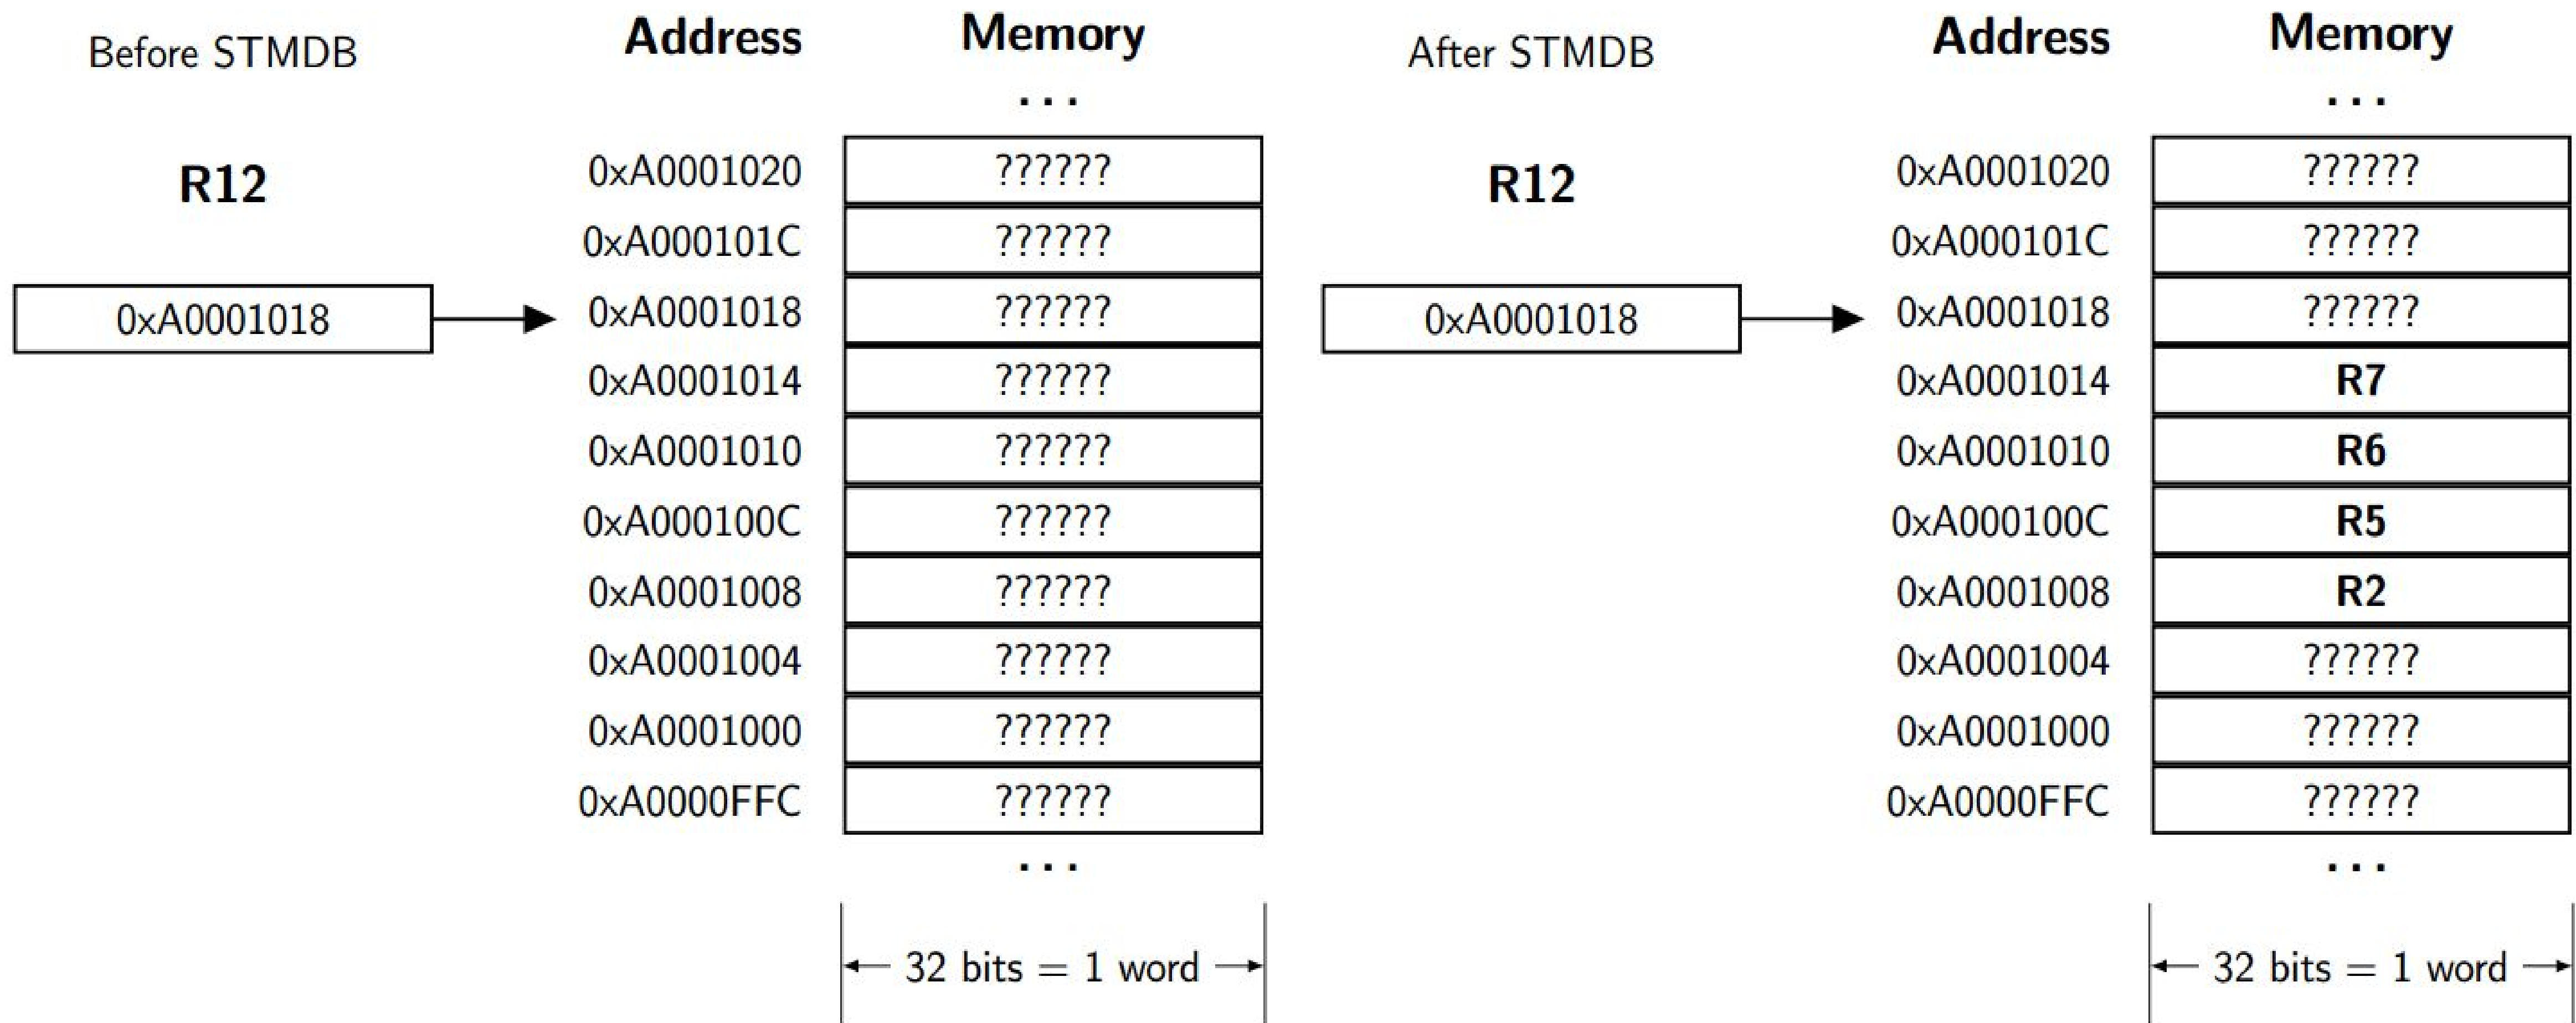
e.g. {R1-R3, R10, R7-R9}

Order in which registers are specified is not important

For both LDM and STM, the lowest register is always loaded from the lowest address, regardless of mode of operation (IA, IB, DA, DB)



STMDB R12, {R5-R7, R2}



...
and now
we're back
to stacks!

LDM and STM instructions can be used to push/pop stack items

Choose mode of operation appropriate to stack growth convention

increment/decrement

before/after

e.g. Full Descending stack

Decrement Before pushing data (STMDB)

Increment After popping data (LDMIA)

To push/pop data using LDM and STM

Use stack pointer register (e.g. R13 or SP) as base register

Use ! syntax to modify LDM/STM behaviour so the stack pointer is updated

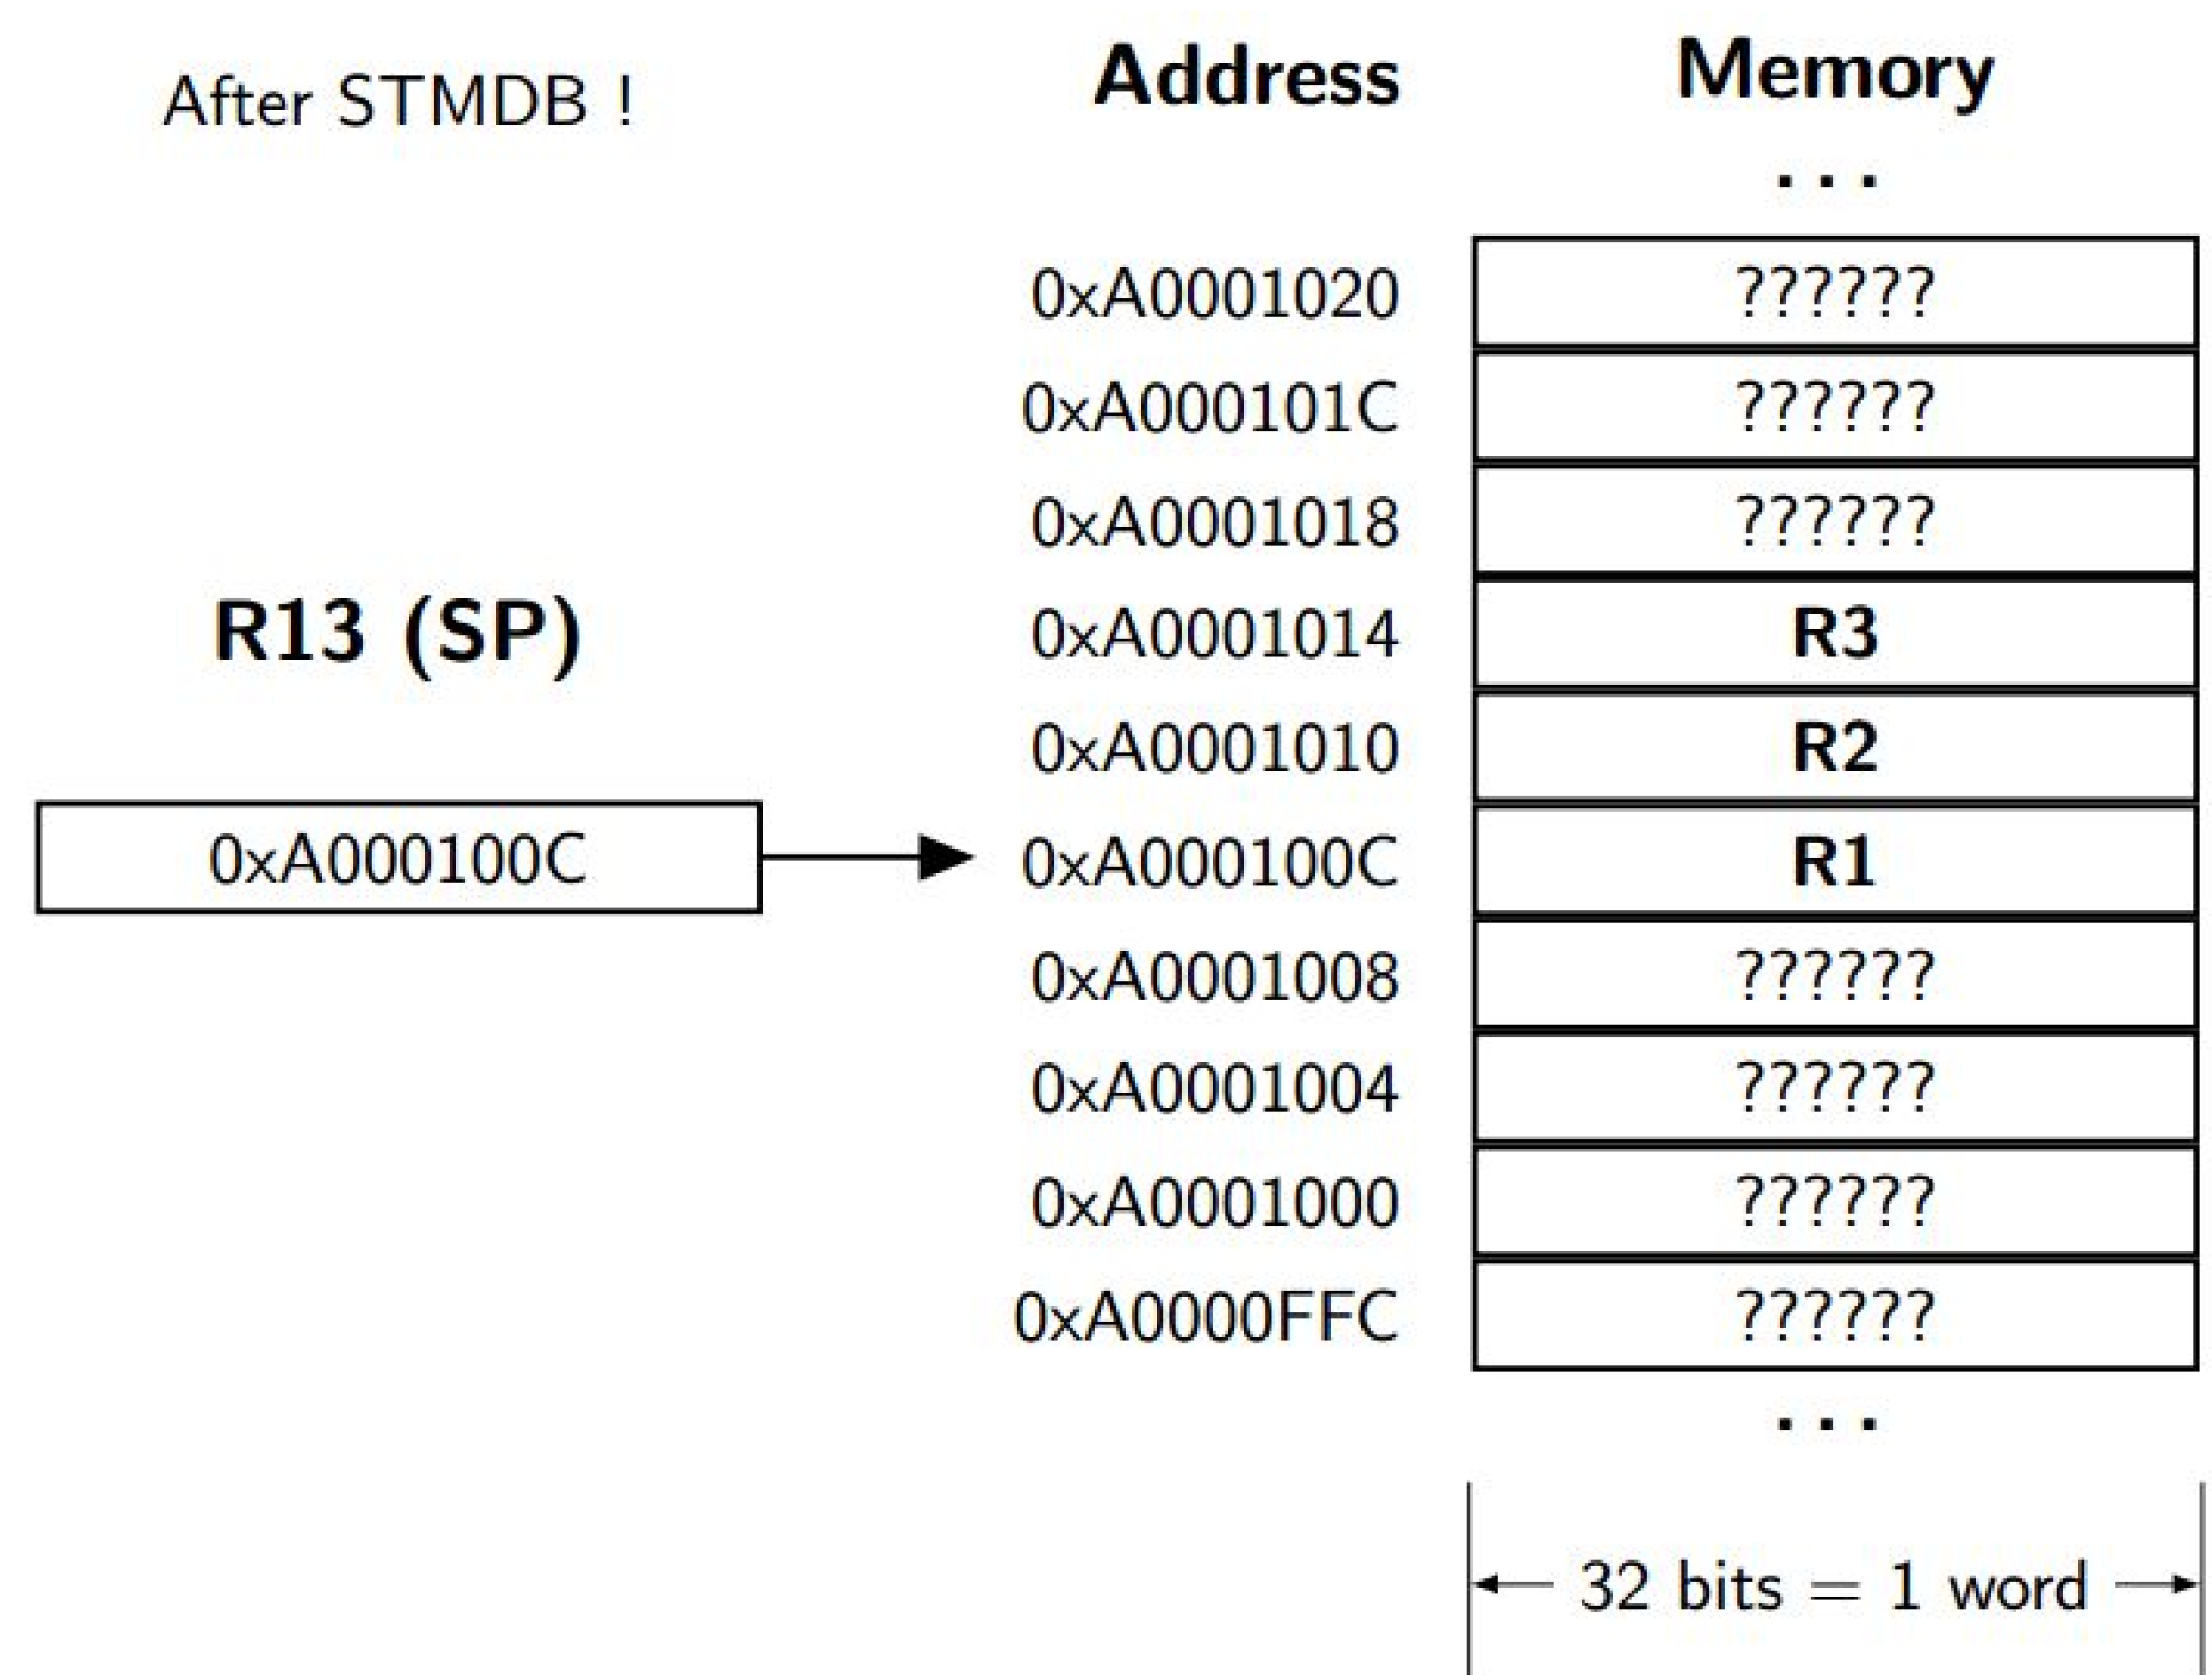
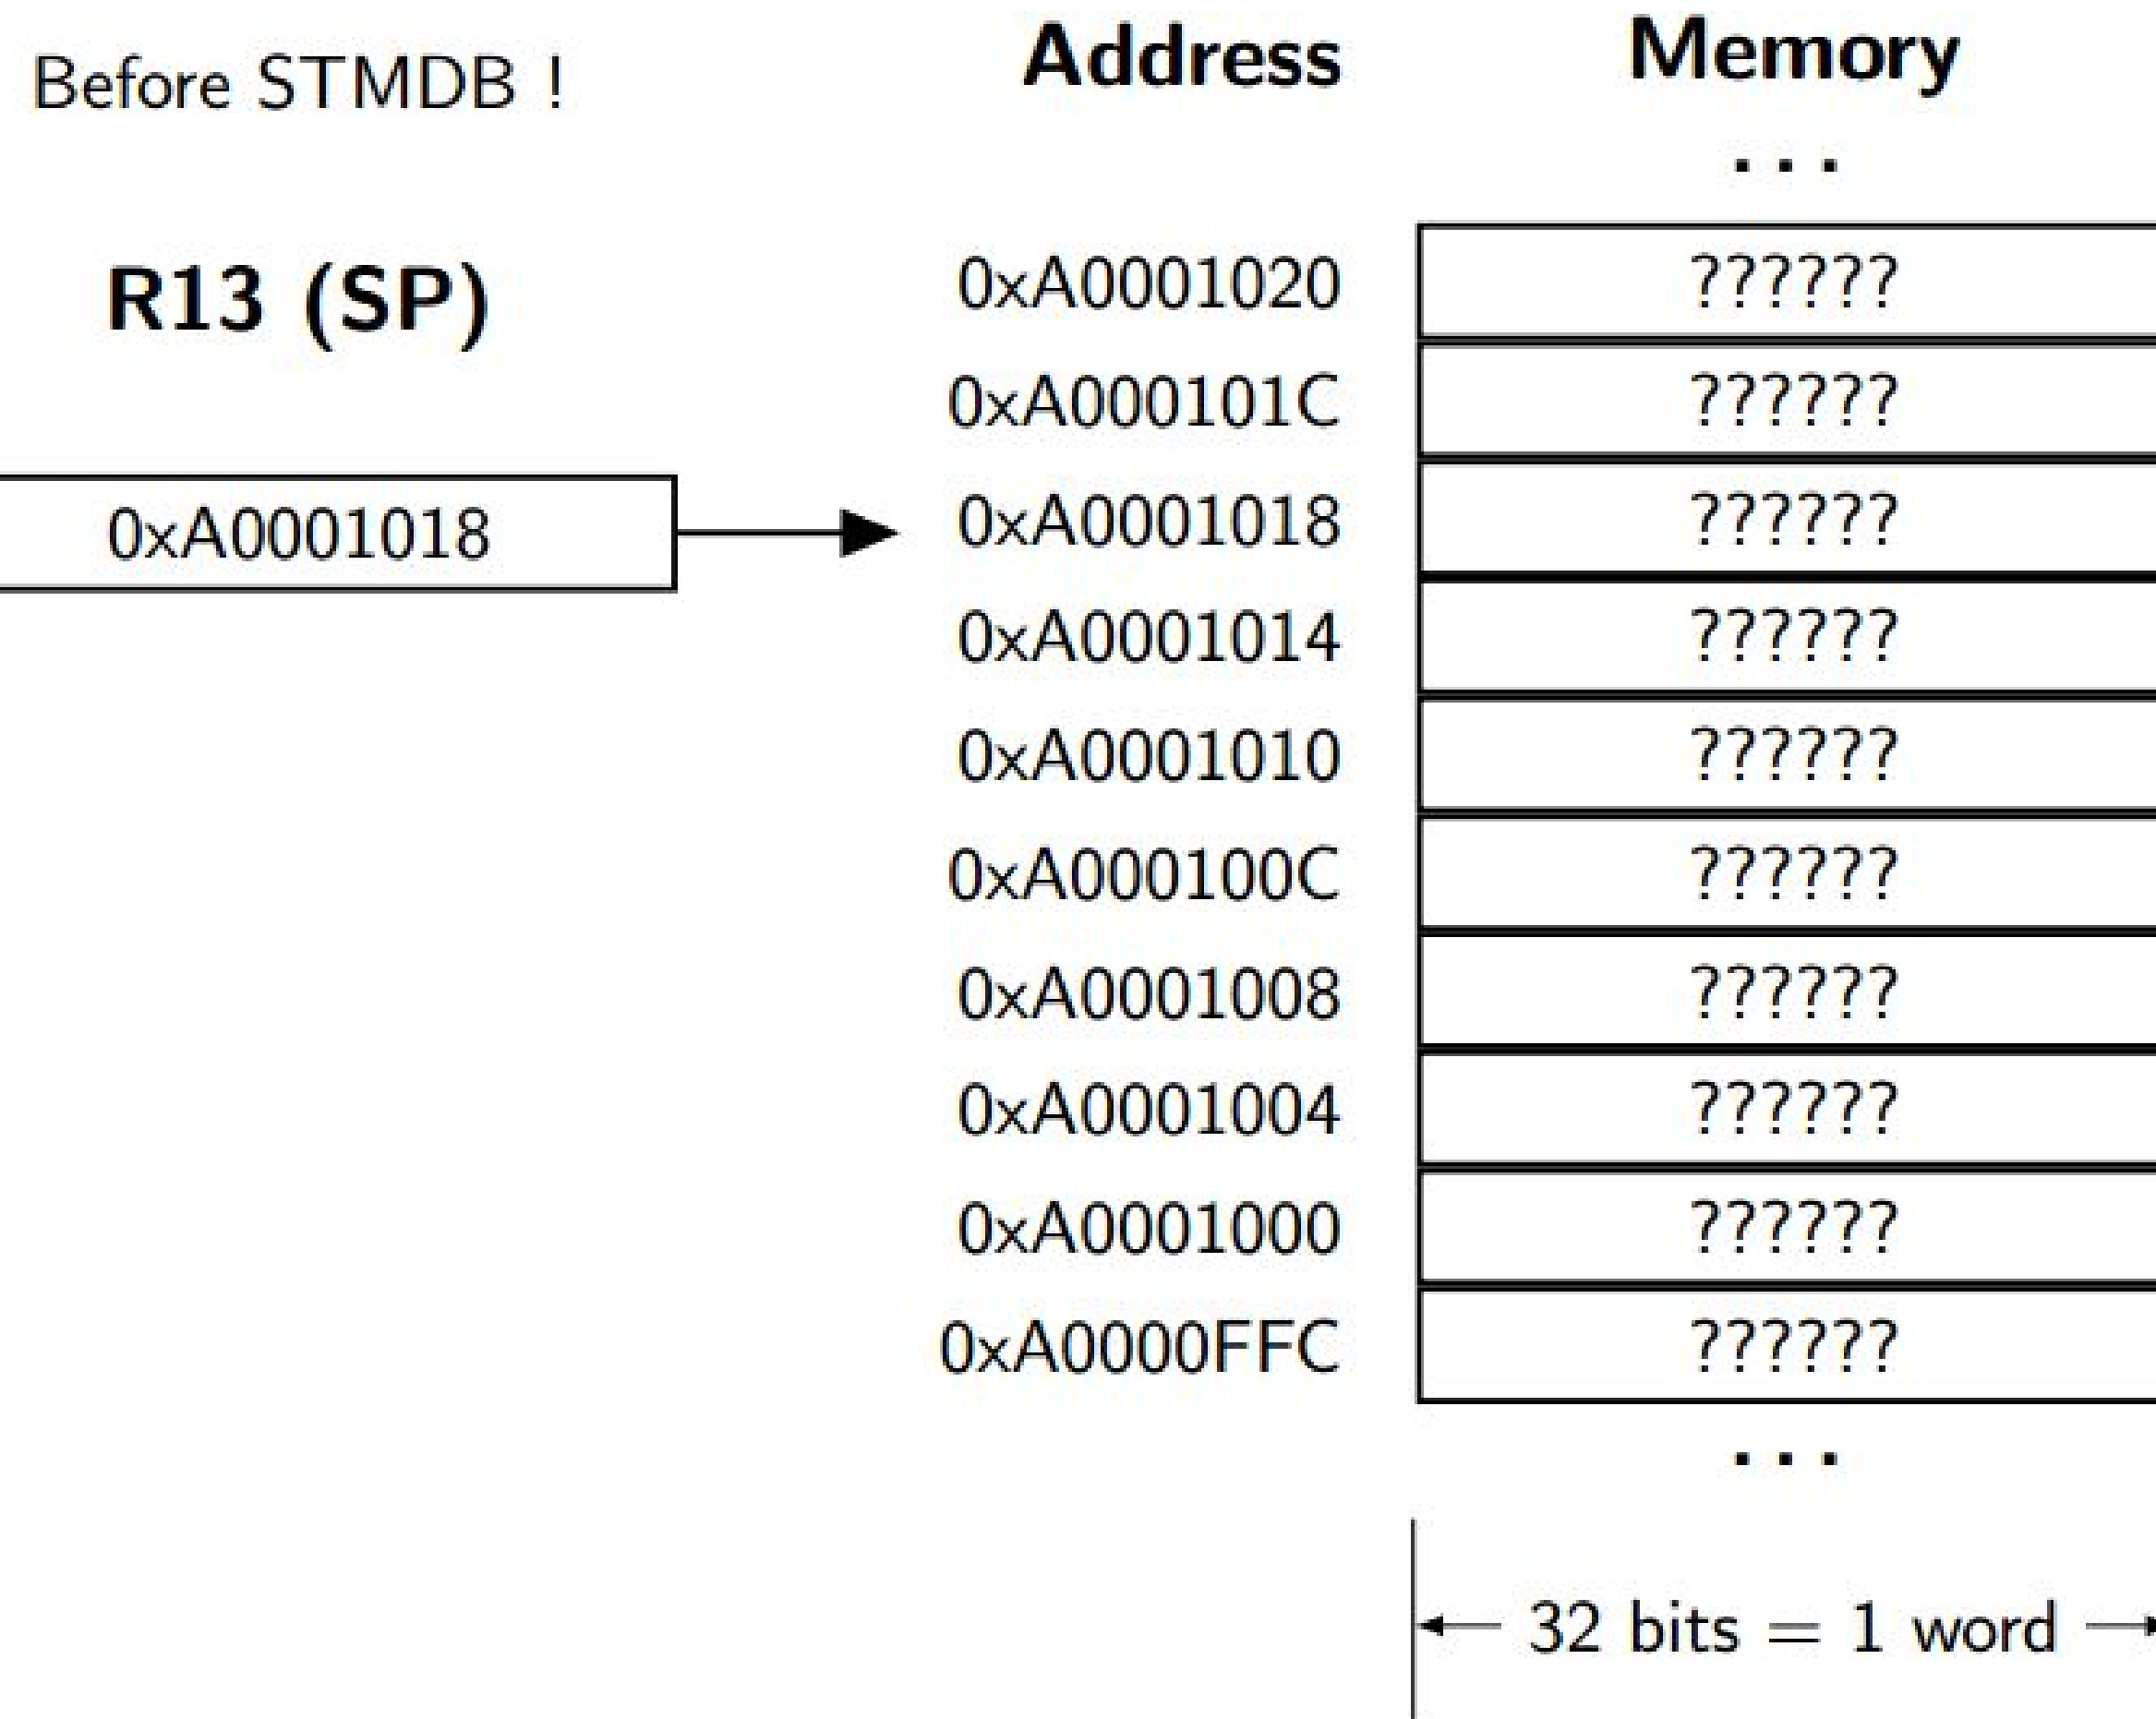
```
STMDB SP!, {R1-R3} ; push R1, R2, R3
```

```
LDMIA SP!, {R1-R3} ; pop R1, R2, R3
```

Note the ! syntax

Push contents of registers R1, R2 and R3

STMDB sp!, {R1-R3}



e.g. Save (push) R1, R2, R3 and R5 on to a full descending stack with R13 (or sp) as the stack pointer

```
STMDB sp!, {R1-R3,R5}
```

Note use of ! in **SP!**

e.g. Restore (pop) R1, R2, R3 and R5 off a full descending stack with R13 (or sp) as the stack pointer

```
LDMIA sp!, {R1-R3,R5}
```

Note use of ! in **SP!**

Works because the lowest register is always loaded from or stored to the lowest address

Stack-oriented synonyms for LDMxx and STMxx allow us to use the same suffix for both LDM and STM instructions

Easier for us to remember!

e.g. Push R1, R2, R3 and R5 on to a full descending stack with R13 (or sp) as the stack pointer

```
STMFD sp!, {R1-R3,R5}
```

e.g. Pop R1, R2, R3 and R5 off a full descending stack with R13 (or sp) as the stack pointer

```
LDMFD sp!, {R1-R3,R5}
```

Stack growth convention	Push		Pop	
	STM mode	stack-oriented synonym	LDM mode	stack-oriented synonym
full descending	STMDB	STMFD	LDMIA	LDMFD
full ascending	STMIB	STMFA	LDMDA	LDMFA
empty descending	STMDA	STMED	LDMIB	LDMED
empty ascending	STMIA	STMEA	LDMDB	LDMEA

Could push values of any size on to a stack

To push a byte from R0 to system stack

```
STRB R0, [SP, #-1] !
```

To pop a byte from system stack to R0

```
LDRB R0, [SP], #1
```

Pushing non-word size data is problematic due to memory alignment constraints

e.g. Push 1 word, followed by 3 half-words, followed by 2 words ...

```
; push word from R0  
STR R0, [SP, #-4]!
```

```
; push 3 half words from R1, R2 and R3  
STRH R1, [SP, #-2]!  
STRH R2, [SP, #-2]!  
STRH R3, [SP, #-2]!
```

```
; push 2 words from R4 and R5  
STR R4, [SP, #-4]!  
STR R5, [SP, #-4]!
```

Won't work as expected

Non-aligned word
access

A stack is a data structure with well defined operations

initialize, push, pop

Stacks are accessed in LIFO order (Last In First Out)

Implemented by

setting aside a region of memory to store the stack contents

initializing a stack pointer to store top-of-stack address

Growth convention

Full/Empty, Ascending/Descending

User defined stack or system stack

When using the system stack, always pop off everything that you push on

not doing this will probably cause an error that may be hard to correct